

Lab 8 - Setting up Model Monitoring

Setup Prometheus and Grafana with HELM

Installing Helm

To install helm version 3 on Linux or MacOS, you can follow following instructions.

```
curl https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3 |  
bash
```

You could further refer to [Official HELM Install Instructions](#) for alternative options.

Verify the installation is successful,

```
helm --help  
helm version
```

Deploy Prometheus Stack with HELM

Read about [kube-prometheus-stack 33.1.0 · prometheus/prometheus-community](#) chart at [artifacthub.io](#)

Add helm repository using ,

```
helm repo add prometheus-community https://prometheus-community.github.io/  
helm-charts  
  
helm repo update
```

Install the helm chart to setup Prometheus and Grafana as,

```
helm upgrade --install prom \
  -n monitoring \
  --create-namespace \
  prometheus-community/kube-prometheus-stack \
  --set grafana.service.type=NodePort \
  --set grafana.service.nodePort=30200 \
  --set prometheus.service.type=NodePort \
  --set prometheus.service.nodePort=30300
```

validate

```
helm list -A
kubectl get all -n monitoring
```

You should be able to access

- prometheus at <http://localhost:30300/>
- grafana at <http://localhost:30200/>

Login to grafana with

```
Username: admin
Password: prom-operator
```

Build Model Monitoring Dashboard

Add Instrumentation for FastAPI

File: `src/api/main.py`

```
from fastapi import FastAPI
from fastapi.middleware.cors import CORSMiddleware
from inference import predict_price, batch_predict
from schemas import HousePredictionRequest, PredictionResponse
from prometheus_fastapi_instrumentator import Instrumentator # 📌 Add this
```

```
# Add CORS middleware
app.add_middleware(
    CORSMiddleware,
    allow_origins=["*"],
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)

# Initialize and instrument Prometheus metrics
Instrumentator().instrument(app).expose(app) # 📌 Add this
```

append this to `requirements.txt`

```
prometheus-fastapi-instrumentator==6.1.0
```

commit and push the changes

```
git add src/api/main.py src/api/requirements.txt
git commit -am "add fastapi instrumentation"
git push origin
```

After pushing the changes, you should see GitHub Actions Pipeline getting triggered ultimately building a new image and publishing it to the DockerHub Repo.

Summary

Jobs

data-processing

model-training

build-and-publish

Run details

Usage

Workflow file

build-and-publish

succeeded 4 minutes ago in 1m 24s

Build and push Docker image

1m 4s

314 #11 exporting manifest list sha256:8c16c301091ddd12a54aef9704f58c962e1160bec2a9aa8cc1693f3fc804ea77 done

315 #11 pushing layers

318 #11 pushing layers 11.6s done

319 #11 pushing manifest for docker.io/initcron/house-price-model:latest@sha256:8c16c301091ddd12a54aef9704f58c962e1160bec2a9aa8cc1693f3fc804ea77

320 #11 pushing manifest for docker.io/initcron/house-price-model:latest@sha256:8c16c301091ddd12a54aef9704f58c962e1160bec2a9aa8cc1693f3fc804ea77 1.6s done

321 #11 DONE 39.8s

322 #13 resolving provenance for metadata file

323 #13 DONE 0.0s

324 ▶ ImageID

326 ▶ Digest

328 ▶ Metadata

370

Post Build and push Docker image

0s

1 Post job cleanup.

2 ▶ Removing temp folder /home/runner/work/_temp/docker-actions-toolkit-km2DVW

3 ▶ Post cache

5

Deploy this change to kubernetes with

```
kubectl rollout restart deployment model
```

If you visit the FastAPI service now, you should see `/metrics` endpoint being added.

House Price Prediction API 1.0.0 OAS3

[/openapi.json](#)

An API for predicting house prices based on various features. This application is part of the MLOps Bootcamp by School of Devops. Authored by Gourav Shah.

[School of Devops - Website](#)

[Send email to School of Devops](#)

Apache 2.0

default		^
GET	/metrics Metrics	
GET	/health Health Check	
POST	/predict Predict	
POST	/batch-predict Batch Predict Endpoint	

Scrape FastAPI Metrics using Prometheus Service Monitor

File : `deployment/monitoring/servicemonitor.yaml`

```
apiVersion: monitoring.coreos.com/v1
kind: ServiceMonitor
metadata:
  name: house-price-api-monitor
  labels:
```

```

    release: prom # Match the label of your Prometheus instance as per helm
release
spec:
  selector:
    matchLabels:
      app: model
  namespaceSelector:
    matchNames:
      - default # or your namespace
  endpoints:
    - port: "8000" # or match name of your service port
      path: /metrics
      interval: 15s

```

apply using

```
kubectl apply -f deployment/monitoring/servicemonitor.yaml
```

check if its been created

```

kubectl get servicemonitor
kubectl describe servicemonitor

```

Validate from Prometheus that the metrics are being sent to prometheus

<http://localhost:30300/targets>

serviceMonitor/default/house-price-api-monitor/0				1/1 up ●	
Endpoint	Labels	Last scrape		State	
http://10.244.1.2:8000/metrics	container="house-price-model" endpoints="http" instances="10.244.1.2:8000" job="model" namespace="default" pod="model-58f75dc4-q16m4" services="model"	1.341s ago		5ms	
				up	

Also try running following queries on prometheus <http://localhost:30300/> if you see the results

```
http_requests_total
```

```

histogram_quantile(0.95, sum(rate(http_request_duration_seconds_bucket[1m]))
by (le, handler))

```

```
rate(http_request_size_bytes_sum[1m])
```

[sample output for http_requests_total]

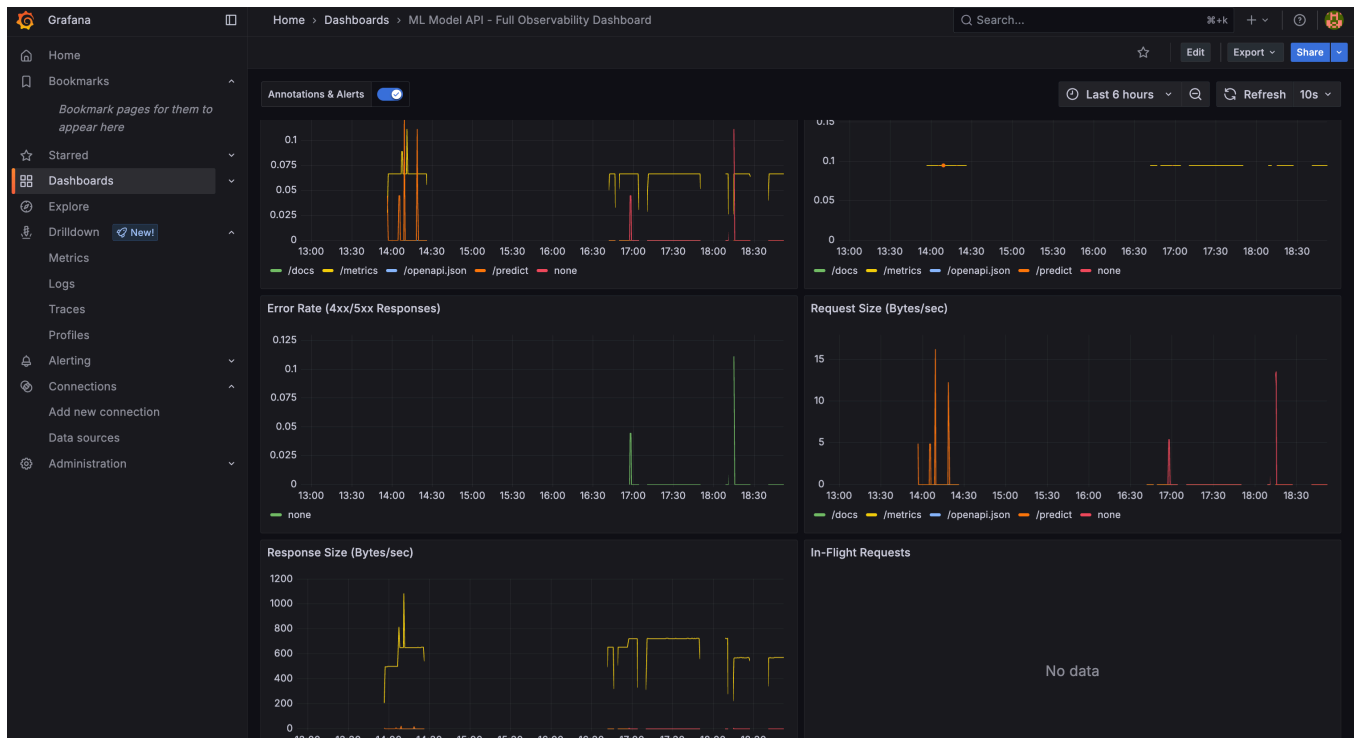
```
http_requests_total{container="house-price-  
model", endpoint="8000", handler="none", instance="10.244.2.8:8000",  
job="model", method="GET", namespace="default", pod="model-6798556657-  
kn6zp", service="model", status="4xx"}  
22  
http_requests_total{container="house-price-model", endpoint="8000", handler="/  
predict", instance="10.244.2.8:8000", job="model", method="POST",  
namespace="default", pod="model-6798556657-  
kn6zp", service="model", status="2xx"}  
2  
http_requests_total{container="house-price-model", endpoint="8000", handler="/  
metrics", instance="10.244.2.8:8000", job="model", method="GET",  
namespace="default", pod="model-6798556657-  
kn6zp", service="model", status="2xx"}  
29
```

If you see the above queries returning some output, it's a confirmation that Prometheus is collecting the data from FastAPI. Now, you just need to visualize it with Grafana.

To do that,

- Login to Grafana
- Go to Dashboards → New → Import
- In the box which reads `Import via dashboard JSON model` paste the code from [enhanced_fastapi_ml_dashboard](#)

You should see a dashboard such as this



This gives you the monitoring which is specific to your model's performance in terms of latency, number of requests, error rates etc. which could be very useful for you to scale your inference later.

#courses/mlops